

# Reinforcement Learning

## 3. Dynamic programming

Olivier Sigaud


Sorbonne Université  
<http://people.isir.upmc.fr/sigaud>



## Introduction


- ▶ Once we have defined an MDP
- ▶ Dynamic programming methods can find the optimal policy
- ▶ Assuming they know everything about the MDP transition and reward function
- ▶ Reinforcement Learning applies when the transition and reward functions are unknown
- ▶ To define dynamic programming methods, we need value functions

## Value and action value functions



- ▶ The **value function**  $V^\pi : S \rightarrow \mathbb{R}$  records the aggregation of reward on the long run for each state (following policy  $\pi$ ). It is a **vector** with one entry per state
- ▶ The **action value function**  $Q^\pi : S \times A \rightarrow \mathbb{R}$  records the aggregation of reward on the long run for doing each action in each state (and then following policy  $\pi$ ). It is a **matrix** with one entry per state and per action
- ▶ In the remainder, we focus on  $V$ , trivial to transpose to  $Q$


## Bellman equation over a Markov chain: recursion



Given the discounted reward aggregation criterion:


- ▶  $V(s_0) = r_0 + \gamma r_1 + \gamma^2 r_2 + \gamma^3 r_3 + \dots$
- ▶  $V(s_0) = r_0 + \gamma(r_1 + \gamma r_2 + \gamma^2 r_3 + \dots)$
- ▶  $V(s_0) = r_0 + \gamma V(s_1)$
- ▶ More generally  $V(s_t) = r_t + \gamma V(s_{t+1})$

## Bellman equation: general case



- ▶ Generalisation of  $V(s_t) = r_t + \gamma V(s_{t+1})$  over all possible trajectories
- ▶ The **expectation** of a random variable is the sum of the realizations weighted by their probabilities
- ▶ The realizations are the next states
- ▶ Deterministic  $\pi$ :  $V^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s'} p(s'|s, \pi(s)) V^\pi(s')$

## Bellman equation: general case



- ▶ Generalisation of  $V(s_t) = r_t + \gamma V(s_{t+1})$  over all possible trajectories
- ▶ The **expectation** of a random variable is the sum of the realizations weighted by their probabilities
- ▶ The realizations are the next states
- ▶ Stochastic  $\pi$ :  $V^\pi(s) = \sum_a \pi(a|s)[r(s, a) + \gamma \sum_{s'} p(s'|s, a)V^\pi(s')]$

## Recursive operators and convergence

- ▶ If we define an operator  $T$  such that  $X_{n+1} \leftarrow TX_n$
- ▶ If  $T$  is **contractive**, then through repeated application of  $T$ ,  $X_n$  will converge to some fixed point
- ▶ For instance, if  $T$  divides by 2,  $X_n$  converges to 0

## The Bellman optimality operator (Value Iteration)

- ▶ We call **Bellman optimality operator** (noted  $T^*$ ) the application

$$V_{n+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s' | s, a) V_n(s') \right]$$

- ▶ If  $\gamma < 1$ ,  $T^*$  is contractive
- ▶ By iterating, computes the value of the current policy
- ▶ The optimal value function is the fixed-point of  $T^*$ :  $V^* = T^* V^*$
- ▶ Value iteration:  $V_{n+1} \leftarrow T^* V_n$



Puterman, M. L. (2014) *Markov decision processes: discrete stochastic dynamic programming*. John Wiley & Sons.



## The Bellman operator (Policy Iteration)

- ▶ We call **Bellman operator** (noted  $T^\pi$ ) the application

$$V_{n+1}^\pi(s) \leftarrow r(s, \pi(s)) + \gamma \sum_{s'} p(s'|s, \pi(s)) V_n^\pi(s')$$

- ▶ If  $\gamma < 1$ ,  $T$  is contractive
- ▶ Converges to optimal value and policy
- ▶ Policy Iteration:

- ▶ Policy evaluation:

$$V_{n+1}^\pi \leftarrow T^\pi V_n^\pi$$

- ▶ Policy improvement:

$$\forall s \in S, \pi'(s) \leftarrow \arg \max_{a \in A} \sum_{s'} p(s'|s, a) [r(s, a) + \gamma V_n^\pi(s')]$$

or

$$\forall s \in S, \pi'(s) \leftarrow \arg \max_{a \in A} [r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_n^\pi(s')]$$

- ▶ Note:  $\sum_{s', r} p(s', r|s, a) [r + \gamma V(s')] = r + \gamma \sum_{s'} p(s'|s, a) V(s')$

## Value Iteration: the algorithm 1st method of dynamic programming

Value Iteration, for estimating  $\pi \approx \pi_*$ 

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation  
 Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$ 
|   Loop for each  $s \in \mathcal{S}$ :
|      $v \leftarrow V(s)$ 
|      $V(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$ 
|      $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
until  $\Delta < \theta$ 
```

Output a deterministic policy,  $\pi \approx \pi_*$ , such that  

$$\pi(s) = \arg \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

- ▶ Taken from Sutton & Barto, 2018, p. 83
- ▶ Reminder:  $\sum_{s',r} p(s', r | s, a) [r + \gamma V(s')] = r + \gamma \sum_{s'} p(s' | s, a) V(s')$

## Value Iteration in practice

0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0		0.0
0.0		0.0		0.0
0.0	0.0	0.0		1

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice

0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0		0.0
0.0		0.0		0.9
0.0	0.0	0.0		1

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice

0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0		0.81
0.0		0.0		0.9
0.0	0.0	0.0		1

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice

0.0	0.0	0.0	0.0	0.73	↖ 0.9*0.81
0.0	0.0	0.0		0.81	↖ 0.9*0.9
0.0		0.0		0.9	
0.0	0.0	0.0		1	→ ?

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice

0.0	0.0	0.0	0.66	0.73
0.0	0.0	0.0		0.81
0.0		0.0		0.9
0.0	0.0	0.0		1

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice

0.0	0.0	0.59	0.66	0.73
0.0	0.0	0.0		0.81
0.0		0.0		0.9
0.0	0.0	0.0		1

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$



## Value Iteration in practice

0.0	0.53	0.59	0.66	0.73
0.0	0.0	0.53		0.81
0.0		0.0		0.9
0.0	0.0	0.0		1


$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice

0.48	0.53	0.59	0.66	0.73
0.43	0.48	0.53		0.81
0.39		0.48		0.9
0.35	0.39	0.43		1

$$\forall s \in S, V_{i+1}(s) \leftarrow \max_{a \in A} \left[ r(s, a) + \gamma \sum_{s'} p(s'|s, a) V_i(s') \right]$$

## Value Iteration in practice



Pourquoi en haut et pas a droite?

We have iterated on values, and determined a policy out of it (without necessarily representing it if using  $Q(s, a)$ )

## Policy Iteration: the algorithm 2nd method for dynamic programming

### Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

- ## 1. Initialization

$V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$

- ## 2. Policy Evaluation

Loop;

$$\Delta \leftarrow 0$$

Loop for each  $s \in S$ :

$$v \leftarrow V(s)$$
$$V(s) \leftarrow \sum_{s',r} p(s',r|s, \pi(s)) [r + \gamma V(s')]$$
$$\Delta \leftarrow \max(\Delta, |v - V(s)|)$$

until  $\Delta < \theta$  (a small positive number determining the accuracy of estimation)

- ### 3. Policy Improvement

*policy-stable*  $\leftarrow$  true

For each  $s \in S$ :

$$old-action \leftarrow \pi(s)$$
$$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$$


If  $old\_action \neq \pi(s)$ , then  $policy\_stable \leftarrow false$

If *policy-stable*, then stop and return  $V \approx v_*$  and  $\pi \approx \pi_*$ ; else go to 2

- Taken from Sutton & Barto, 2018, p. 80


- Note:  $\sum_{s', r} p(s', r|s, a)[r + \gamma V(s')] = r + \gamma \sum_{s'} p(s'|s, a)V(s')$

## Policy Iteration in practice




$$\forall s \in S, V_i(s) \leftarrow \text{evaluate}(\pi_i(s))$$

## Policy Iteration in practice




$$\forall s \in S, \pi_{i+1}(s) \leftarrow \text{improve}(\pi_i(s), V_i(s))$$

## Policy Iteration in practice



$$\forall s \in S, V_i(s) \leftarrow \text{evaluate}(\pi_i(s))$$


## Policy Iteration in practice



$$\forall s \in S, \pi_{i+1}(s) \leftarrow \text{improve}(\pi_i(s), V_i(s))$$




## Policy Iteration in practice




$$\forall s \in S, V_i(s) \leftarrow \text{evaluate}(\pi_i(s))$$

## Policy Iteration in practice




$$\forall s \in S, \pi_{i+1}(s) \leftarrow \text{improve}(\pi_i(s), V_i(s))$$

## Policy Iteration in practice




$$\forall s \in S, V_i(s) \leftarrow \text{evaluate}(\pi_i(s))$$

## Policy Iteration in practice




$$\forall s \in S, \pi_{i+1}(s) \leftarrow \text{improve}(\pi_i(s), V_i(s))$$

## Policy Iteration in practice




$$\forall s \in S, V_i(s) \leftarrow \text{evaluate}(\pi_i(s))$$

## Policy Iteration in practice




$$\forall s \in S, \pi_{i+1}(s) \leftarrow \text{improve}(\pi_i(s), V_i(s))$$

## Policy Iteration in practice



$$\forall s \in S, V_i(s) \leftarrow \text{evaluate}(\pi_i(s))$$



## Policy Iteration in practice



Here we have managed a policy and a value representations at all steps



## Generalized Policy Iteration



- Policy iteration evaluates each intermediate policy up to convergence.  
**This is slow.**
- Instead, evaluate the policy for  $N$  iterations, or even not for all states.
- **Asynchronous dynamics programming**: decoupling policy evaluation and improvement
- Taken from Sutton & Barto, 2018

## Corresponding labs

- ▶ Implement value iteration with the  $V$  and the  $Q$  functions
- ▶ Implement policy iteration with the  $V$  and the  $Q$  functions
- ▶ Compare them
- ▶ Optional: study “Generalized policy iteration”

Any question?



Send mail to: [Olivier.Sigaud@upmc.fr](mailto:Olivier.Sigaud@upmc.fr)



Puterman, M. L.

*Markov decision processes: discrete stochastic dynamic programming.*

John Wiley & Sons, 2014.