

# Reinforcement Learning

## 5. Off-policy versus on-policy RL

Olivier Sigaud


Sorbonne Université  
<http://people.isir.upmc.fr/sigaud>



# Introduction

- ▶ We have said that SARSA was **on-policy** and Q-LEARNING was **off-policy**.
- ▶ The goal of this class is to understand precisely what this means
- ▶ This distinction matters a lot for deep RL research
- ▶ Off-policy algorithms like DDPG are generally more sample efficient but less stable than on-policy algorithms like PPO

## Basic concepts



- ▶ To understand the distinction, one must consider three objects:
  - ▶ The behavior policy  $\beta(s)$  used to generate samples.
  - ▶ The critic, which is generally  $V(s)$  or  $Q(s, a)$
  - ▶ The target policy  $\pi(s)$  used to control the system in exploitation mode.

## Off-policiness: definition

- ▶ “Off-policy learning” refers to learning about one way of behaving, called the *target policy*, from data generated by another way of selecting actions, called the *behavior policy*.
- ▶ Two notions:
  - ▶ Off-policy policy evaluation
  - ▶ Off-policy control



Maei, H. R., Szepesvári, C., Bhatnagar, S., & Sutton, R. S. (2010) Toward off-policy learning control with function approximation. *ICML*, pages 719–726.

# Off-policy policy evaluation: Definition

0.26	0.23	0.59	0.66	0.73
0.29	0.21	0.53		0.81
0.32		0.48		0.9
0.35	0.39	0.43		1

 $\beta(s)$ 

0.48	0.53	0.59	0.66	0.73
0.43	0.48	0.53		0.81
0.39		0.48		0.9
0.35	0.39	0.43		1

 $\pi(s)$ 

- ▶ Can evaluate the critic of a target policy  $\pi(s)$  from playing a different behavior policy  $\beta(s)$ ?
- ▶ Obviously,  $\beta(s)$  and  $\pi(s)$  generate different values  $V(s)$  or  $Q(s, a)$
- ▶ The goal of “off-policy correction” is to correct for the sample mismatch
- ▶ The target policy does not need to be optimal
- ▶ This is a weak notion of off-policiness (not covered here)



Precup, D. (2000) Eligibility traces for off-policy policy evaluation. *Computer Science Department Faculty Publication Series*



Munos, R., Stepleton, T., Harutyunyan, A., & Bellemare, M. G. (2016) Safe and efficient off-policy reinforcement learning. In *Advances in Neural Information Processing Systems*, pages 1054–1062


## Off-policy control: Definition

- ▶ Whatever the behavior policy (as few assumptions as possible)
- ▶ The target policy should be an approximation to the optimal policy
- ▶ Ex: stochastic behavior policy, deterministic target policy
- ▶ An algorithm might be more or less off-policy depending on the assumptions on  $\beta(s)$

## Why preferring off-policy to on-policy control?

- ▶ More freedom for exploration
- ▶ Learning from human data (imitation)
- ▶ Reusing old data, e.g. from a replay buffer (sample efficiency)
- ▶ Transfer between policies in a multitask context


## Approach



- ▶ Two steps: open-loop study then closed-loop study
  - ▶ Use uniform sampling as “behavior policy” (few assumptions)
  - ▶ No exploration issue, no bias towards good samples
  - ▶ NB: in uniform sampling, samples do not correspond to an agent trajectory
  - ▶ An alternative is random walk, but may raise exploration issues
  - ▶ Study critic learning from these samples
- ▶ Then close the loop:
  - ▶ Use the target policy + some exploration as behavior policy
  - ▶ If the target policy gets good, bias more towards good samples



## Learning a critic from samples





Two random actions

- ▶ General format of samples  $S$ :  $(s_t, a_t, r_t, s_{t+1}, a')$
- ▶ Makes it possible to apply a general update rule:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_t + \gamma Q(s_{t+1}, a') - Q(s_t, a_t)]$$


- ▶ There are three possible update rules:
  1.  $a' = \operatorname{argmax}_a Q(s_{t+1}, a)$  (corresponds to Q-LEARNING)
  2.  $a' = \beta(s_{t+1})$  (corresponds to SARSA)
  3.  $a' = \pi(s_{t+1})$  (corresponds e.g. to DDPG, an ACTOR-CRITIC algorithm)

## Off-policiness with uniform sampling




- We add a negative reward for hitting walls

## Detailed mechanism: Rule 1 (Q-LEARNING)




- ▶  $a' = \operatorname{argmax}_a Q(s_{t+1}, a)$
- ▶ Always backpropagates the highest value
- ▶ Thus  $Q(s, a)$  consistently converges to the value of acting optimally

## Detailed mechanism: Rule 2 (SARSA)




- ▶  $a' = \beta(s_{t+1})$
- ▶ Due to uniform sampling, the probabilities to take any  $a'$  are uniform
- ▶ Thus  $Q(s, a)$  converges to some average, corresponding to performing random walk
- ▶ Being greedy wrt  $Q(s, a)$  does not result in the optimal target policy

## Detailed mechanism: Rule 3 (DDPG-like)





- ▶  $a' = \pi(s_{t+1})$
- ▶  $\pi(s_{t+1})$  evolve consistently with action values (update of  $\alpha\delta$ )
- ▶ Thus actions with higher values get sampled more often (vs uniformly in SARSA)
- ▶ If a good  $a'$  is sampled, increase the updated value
- ▶ If a bad  $a'$  is sampled, decrease it
- ▶ Results in more structured fluctuations

## Results



- ▶ Rule 1 learns an optimal critic (thus Q-LEARNING is truly off-policy)
- ▶ Rule 2 fails (thus SARSA is not off-policy)
- ▶ Rule 3 fails too (thus an algorithm like DDPG is not truly off-policy!)
- ▶ NB: different ACTOR-CRITIC implementations behave differently
- ▶ E.g. if the critic estimates  $V(s)$ , then equivalent to Rule 1

## Closing the loop Sarsa and DDPG will only work when $\beta(s)=\pi^*(s)$ !!!!!!



- ▶ If  $\beta(s) = \pi^*(s)$ , then Rules 2 and 3 are equivalent,
- ▶ Furthermore,  $Q(s, a)$  will converge to  $Q^*(s, a)$ , and Rule 1 will be equivalent too.
- ▶ Quite obviously, Q-LEARNING still works
- ▶ SARSA and ACTOR-CRITIC work too: the  $\beta(s)$  becomes “Greedy in the limit of infinite exploration” (GLIE)



Singh, S. P., Jaakkola, T., Littman, M. L., & Szepesvári, C. (2000) Convergence results for single-step on-policy reinforcement-learning algorithms. *Machine learning*, 38(3):287–308

## Corresponding labs

- ▶ See [https://github.com/osigaud/rl\\_labs\\_notebooks](https://github.com/osigaud/rl_labs_notebooks)
- ▶ One notebook about off-policy versus on-policy learning
- ▶ Check the three rules convergence properties when  $\beta(s)$  is uniform random sampling



Any question?



Send mail to: [Olivier.Sigaud@upmc.fr](mailto:Olivier.Sigaud@upmc.fr)



Maei, H. R., Szepesvári, C., Bhatnagar, S., & Sutton, R. S. (2010).

Toward off-policy learning control with function approximation.

Édité dans *ICML*, pages 719–726.



Munos, R., Stepleton, T., Harutyunyan, A., & Bellemare, M. G. (2016).

Safe and efficient off-policy reinforcement learning.

Édité dans *Advances in Neural Information Processing Systems*, pages 1054–1062.



Precup, D. (2000).

Eligibility traces for off-policy policy evaluation.

*Computer Science Department Faculty Publication Series*, page 80.



Singh, S. P., Jaakkola, T., Littman, M. L., & Szepesvári, C. (2000).

Convergence results for single-step on-policy reinforcement-learning algorithms.

*Machine learning*, 38(3):287–308.