

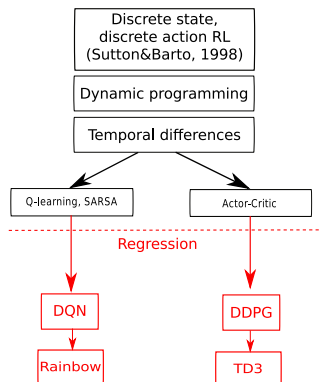
(Deep) Policy Search and Policy Gradient

Olivier Sigaud

Sorbonne Université
<http://people.isir.upmc.fr/sigaud>



Standard RL Class overview

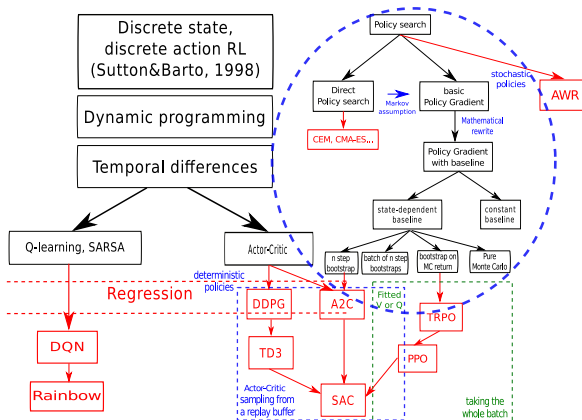


► From Sutton&Barto to deep RL...



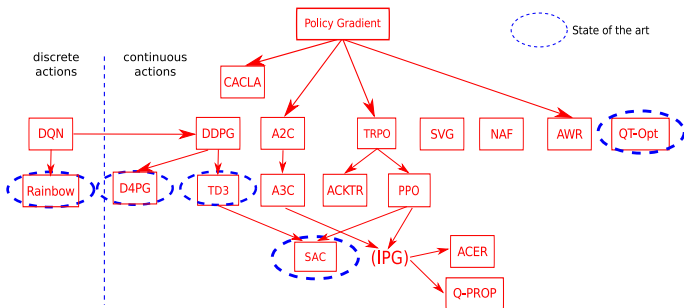
Sutton, R. S. & Barto, A. G. (1998) *Reinforcement Learning: An Introduction*. MIT Press.

Deep Policy Search overview



- ▶ Builds on “Deep RL bootcamp” youtube videos
https://www.youtube.com/watch?v=S_gwYj1Q-44
- ▶ Differences between “pure” policy gradient and actor critic

Next video



- Overview of the most important state-of-the-art deep policy search algorithms
- Main concepts and properties
- Plus videos for individual algorithms

General Goal of Policy Search



► Let:

- π_θ be the parametrized policy of an agent
- τ_θ is an agent trajectory
- $R(\tau_\theta)$ is the corresponding return
- $J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]$ is the global utility (or cost) function

- We have to sample the expectation, thus the goal is to find


$$\theta^* = \operatorname{argmax}_{\theta} J(\theta) = \operatorname{argmax}_{\theta} \sum_{\tau} P(\tau, \theta) R(\tau) \quad (1)$$

- where $P(\tau, \theta)$ is the probability of τ under policy π_θ
- We are in a black-box context: we choose a θ , we generate trajectories and get the return $J(\theta)$ of these trajectories
- Then we look for a better θ



Deisenroth, M. P., Neumann, G., Peters, J., et al. (2013) A survey on policy search for robotics. *Foundations and Trends® in Robotics*, 2(1–2):1–142

(Truly) Random Search



- ▶ Select θ_i randomly
- ▶ Perform a set of τ and get $\hat{J}(\theta_i)$
- ▶ If $\hat{J}(\theta_i)$ is the best so far, keep θ_i
- ▶ Loop until $\hat{J}(\theta_i) > target$


- ▶ Of course, this is not efficient if the space of θ is large
- ▶ General “blind” algorithm, no assumption on $J(\theta)$
- ▶ We can do better if $J(\theta)$ shows some local regularity




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ Start with a policy π_θ with performance $J(\theta)$
- ▶ Generate random variations of π_{θ_i} and evaluate their performance $J(\theta_i)$




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ Select the best variations, ignore the rest
- ▶ Get a new policy π_{θ} from selected variations




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ Repeat the same process
- ▶ Approximates the gradient




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ If variations are wide enough, may escape from easy local minima
- ▶ Covariance matrices adapt the width of variations




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ If variations are wide enough, may escape from easy local minima
- ▶ Covariance matrices adapt the width of variations




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ Until stuck into a wide local minimum
- ▶ Genetic Algorithms, Evolution Strategies, Finite Differences...




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- Compute the local derivative
- Provides steepest descent




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ Follow the gradient with a step
- ▶ Necessity to tune step size




Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search



Gradient descent




- ▶ Iterate until no more improvement
- ▶ Stochastic variant escapes too local minima



Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Two families of methods

Direct Policy Search




- Needs many samples
- More easily escapes local minima
- A separate class about this topic



Sigaud, O. & Stulp, F. (2019) Policy search in continuous action domains: an overview. *Neural Networks*, 113:28-40

Gradient descent



- No sample needed
- Gets stuck into local minima
- $J(\theta)$ unknown in policy search
- Solution: policy gradient methods

Policy gradient methods

- ▶ Direct policy search uses $\langle \theta, J(\theta) \rangle$ pairs and directly looks for θ with the highest $J(\theta)$
- ▶ It ignores the fact that the return comes from state and action trajectories generated by a controller π_θ
- ▶ We can use explicit gradients if we take this information into account
 - ▶ Represent a family of stochastic policies
 - ▶ Increase the probabilities of actions producing trajectories with a high return
 - ▶ Not black-box anymore: access the state, action and immediate reward at each step
 - ▶ The transition and reward functions are still unknown (gray-box approach)
- ▶ Watch Pieter Abbeel's deep RL bootcamp video #4A:
https://www.youtube.com/watch?v=S_gwYj1Q-44

- ▶ Reminder: we look for $\theta^* = \operatorname{argmax}_{\theta} J(\theta) = \operatorname{argmax}_{\theta} \sum_{\tau} P(\tau, \theta) R(\tau)$

$$= \sum_{\tau} \nabla_{\theta} P(\tau, \theta) R(\tau)$$

* gradient of sum is sum of gradients

$$= \sum_{\tau} \frac{P(\tau, \theta)}{P(\tau, \theta)} \nabla_{\theta} P(\tau, \theta) R(\tau)$$

* Multiply by one

$$= \sum_{\tau} P(\tau, \theta) \frac{\nabla_{\theta} P(\tau, \theta)}{P(\tau, \theta)} R(\tau)$$

* Move one term

$$= \sum_{\tau} P(\tau, \theta) \nabla_{\theta} \log P(\tau, \theta) R(\tau)$$

* by property of gradient of log

$$= \mathbb{E}_{\tau}[\nabla_{\theta} \log P(\tau, \theta) R(\tau)]$$

* by definition of the expectation

- ▶ The expectation can be approximated over m trajectories

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}, \theta) R(\tau^{(i)})$$

https://www.youtube.com/watch?v=S_gwYj1Q-44

(12')

Plain Policy Gradient (step 2)

- ▶ We do not have an analytical expression for $P(\tau, \theta)$
- ▶ Thus the gradient $\nabla_{\theta} \log P(\tau^{(i)}, \theta) R(\tau^{(i)})$ cannot be computed
- ▶ Let us reformulate $P(\tau, \theta)$ using the policy π_{θ}

$$P(\tau^{(i)}, \theta) = \prod_{t=1}^H p(s_{t+1}^{(i)} | s_t^{(i)}, a_t^{(i)}) \cdot \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \quad (2)$$

- ▶ (Strong) Markov assumption here
- ▶ At each step, probability of taking each action (defined by the policy) times probability of reaching the next state given the action
- ▶ Then product over states for the whole horizon H

https://www.youtube.com/watch?v=S_gwYj1Q-44 (18')

Plain Policy Gradient (step 2 continued)

► Thus

$$\begin{aligned}\nabla_{\theta} \log P(\tau^{(i)}, \theta) &= \nabla_{\theta} \log \left[\prod_{t=1}^H p(s_{t+1}^{(i)} | s_t^{(i)}, a_t^{(i)}) \cdot \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \right] \\ &\quad * \text{log of product is sum of logs} \tag{3} \\ &= \nabla_{\theta} \left[\sum_{t=1}^H \log p(s_{t+1}^{(i)} | s_t^{(i)}, a_t^{(i)}) + \sum_{t=1}^H \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \right] \\ &= \nabla_{\theta} \sum_{t=1}^H \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \quad * \text{because first term independent of } \theta \\ &= \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \quad * \text{no dynamics model required!}\end{aligned}$$

https://www.youtube.com/watch?v=S_gwYj1Q-44 (18')

Plain Policy Gradient (step 2 continued)

- Reminder

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}, \theta) R(\tau^{(i)})$$


- Thus

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) R(\tau^{(i)}) \quad (5)$$

- The policy structure π_{θ} is known, thus the gradient $\nabla_{\theta} \log \pi_{\theta}$ can be computed
- Can be turned into a practical (but inefficient) algorithm
- We moved from direct policy search on $J(\theta)$ to gradient descent on π_{θ}


https://www.youtube.com/watch?v=S_gwYj1Q-44 (18')

Practical algorithm 1: overview



- ▶ Collect a set of m trajectories $(s_t^{(i)}, a_t^{(i)}, r_t^{(i)}), i \in \{1, H\}$
- ▶ Compute the resulting return $R(\tau^{(i)}) = \sum_{t=1}^H r_t^{(i)}$.
- ▶ For each visited $(s_t^{(i)}, a_t^{(i)})$ pair, apply $\nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \cdot R(\tau^{(i)})$
- ▶ Given (5), this ensures $J(\theta)$ will improve
- ▶ Loop until $J(\theta)$ reaches a local optimum or after some budget


Practical algorithm 1: intuition



$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) R(\tau^{(i)})$$

- ▶ PG is quite different from DPS: search in the state-action space versus parameter space, using some structural assumptions
- ▶ Increasing the log proba. of rewarded actions taken in states increases $J(\theta)$
- ▶ $R(\tau)$ is the step size of each gradient update
- ▶ A bigger $R(\tau)$ results in a bigger update


Practical algorithm 1: further intuition



- Probabilities π_θ must sum to 1, thus increasing one decreases the others
- Moves the action probabilities π_θ in each state towards those providing the highest $R(\tau)$


https://www.youtube.com/watch?v=S_gwYj1Q-44

Distributions over actions: Bernoulli




- ▶ Binary choice between two actions
- ▶ p is a probability, must keep between 0 and 1
- ▶ Use sigmoid, or tanh...
- ▶ Increasing $p(\text{left})$ decreases $p(\text{right})$

Distributions over actions: Categorical





- Choice between K discrete actions
- All the probabilities must sum to 1
- When increasing one probability, how should we decrease the others?
(renormalize?)

Distributions over actions: Normal



- ▶ Choice of a continuous action (extension to multidimensional with multivariate Gaussian)
- ▶ The integral must keep to 1
- ▶ Standard approach: keep variance σ constant

Policy representation (continuous action)



The stochastic policy is represented as a multivariate Gaussian:

$$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) = \mathcal{N}(\boldsymbol{\mu}_{\theta}, \boldsymbol{\Sigma}) = e^{-\frac{1}{2}(\boldsymbol{\mu}_{\theta} - \mathbf{a}_t)^T \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_{\theta} - \mathbf{a}_t)}$$

$$\log \pi_{\theta}(a_t | s_t) = -\frac{1}{2}(\boldsymbol{\mu}_{\theta} - \mathbf{a}_t)^T \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_{\theta} - \mathbf{a}_t)$$



$$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) = \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_{\theta} - \mathbf{a}_t)$$

Just backpropagate $\boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_{\theta} - \mathbf{a}_t) \sum_t (r(s_t, \mathbf{a}_t))$

- NB: We considered a fixed $\boldsymbol{\Sigma}$.
- Learning $\boldsymbol{\Sigma}_{\theta}$ results in a more involved derivation (but provided by librairies)

<https://www.youtube.com/watch?v=SQtOI9jsrJO>

Policy representation (1D continuous action case)



The stochastic policy is represented as a Gaussian:

$$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) = e^{-\frac{1}{2} \frac{(\mu_{\theta} - a_t)^2}{\sigma}}$$

$$\log \pi_{\theta}(a_t | s_t) = -\frac{1}{2} \frac{(\mu_{\theta} - a_t)^2}{\sigma}$$

$$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) = -\frac{\mu_{\theta} - a_t}{\sigma}$$

Just backpropagate $-\frac{\mu_{\theta} - a_t}{\sigma} \sum_t r(s_t, a_t)$

<https://www.youtube.com/watch?v=SQt0I9jsrJ0>

Back to Plain Policy Gradient (step 3)

- ▶ Algo. 1 takes a large batch of trajectories: suffers from large variance
- ▶ Computing from complete trajectories is not the best we can do


$$\begin{aligned}\nabla_{\theta} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) R(\tau^{(i)}) \\ &= \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \left[\sum_{t=1}^H r(s_t^{(i)}, a_t^{(i)}) \right] \\ &\quad * \text{split into two parts} \tag{6}\end{aligned}$$

$$\begin{aligned}&= \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \left[\sum_{k=0}^{t-1} r(s_k^{(i)}, a_k^{(i)}) + \sum_{k=t}^H r(s_k^{(i)}, a_k^{(i)}) \right] \\ &\quad * \text{past rewards do not depend on the future} \tag{7}\end{aligned}$$

$$= \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \left[\sum_{k=t}^H r(s_k^{(i)}, a_k^{(i)}) \right]$$

https://www.youtube.com/watch?v=S_gwYj1Q-44 (26')

Algorithm 2: from step 2 to step 3



- Same as Algorithm 1
- But computes the sum backwards
- Slightly better algorithm

Plain Policy Gradient (step 3 continued)

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \left[\sum_{k=t}^H r(s_k^{(i)}, a_k^{(i)}) \right] \quad (8)$$

- We can reduce the variance by discounting the rewards along the trajectory

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \left[\sum_{k=t}^H \gamma^k r(s_k^{(i)}, a_k^{(i)}) \right]$$

►


$$\sum_{k=t}^H \gamma^k r(s_k^{(i)}, a_k^{(i)}) \text{ can be rewritten } Q^{\pi}(s_t^{(i)}, a_t^{(i)})$$

- Thus we get

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) Q^{\pi}(s_t^{(i)}, a_t^{(i)})$$

https://www.youtube.com/watch?v=S_gwYj1Q-44 (26')

Algorithm 3: discounting the reward



- Q^π is estimated from Monte Carlo
- Even smaller variance

Policy Gradient with constant baseline


- ▶ Besides, we can subtract a “baseline” to (8) without changing its mean, but improving its variance

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)}) \left[\sum_{k=t}^H r(s_k^{(i)}, a_k^{(i)}) - b(s_t^{(i)}) \right]$$

- ▶ A first baseline is the average return \bar{r} over all states of the batch
- ▶ We then normalize each local return with $r_t^{(i)} - \bar{r}$ and divide by the standard deviation so as to get a mean of 0 and a standard deviation of 1.
- ▶ Greater than average returns get positive, smaller get negative
- ▶ Suggested in <https://www.youtube.com/watch?v=tqrcjHuNdmQ>

Algorithm 4: adding a constant baseline

- ▶ Estimate \bar{r} and $std(r)$ from all rollouts
- ▶ Same as Algorithm 2, using $(r_t^{(i)} - \bar{r})/std(r)$



- ▶ Suffers from even less variance

Policy Gradient with state-dependent baseline

- ▶ A better baseline is
$$b(s_t) = \mathbb{E}_\tau[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{H-t} r_H] = V^\pi(s_t)$$
- ▶ The expectation can be approximated from the batch of trajectories
- ▶ Thus we get

$$\nabla_\theta J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_\theta \log \pi_\theta(a_t^{(i)} | s_t^{(i)}) [Q^{\pi_\theta}(s_t^{(i)} | a_t^{(i)}) - V^\pi(s_t^{(i)})]$$

- ▶ $A^\pi(s_t, a_t) = Q^\pi(s_t | a_t) - V^\pi(s_t)$ is the advantage function
- ▶ And we get

$$\nabla_\theta J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H \nabla_\theta \log \pi_\theta(a_t^{(i)} | s_t^{(i)}) A^\pi(s_t^{(i)}, a_t^{(i)})$$

https://www.youtube.com/watch?v=S_gwYj1Q-44 (27')



Williams, R. J. (1992) Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3-4):229–256



Algorithm 5: adding a state-dependent baseline

- ▶ Estimate $V^\pi(s_t)$ from all rollouts
- ▶ Estimate $A^\pi(s_t^{(i)}|a_t^{(i)})$ from all rollouts
- ▶ Same as Algorithm 1 with $A^\pi(s_t^{(i)}|a_t^{(i)})$ instead of $R(\tau^{(i)})$
- ▶ Suffers from even less variance
- ▶ Still no bootstrap update of an estimate \hat{V}_ϕ or \hat{Q}_ϕ

State-dependent baseline: towards bootstrap

Algorithm 1 “Vanilla” policy gradient algorithm

Initialize policy parameter θ , baseline b

for iteration=1, 2, ... **do**

 Collect a set of trajectories by executing the current policy

 At each timestep in each trajectory, compute

 the *return* $R_t = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'}$, and

 the *advantage estimate* $\hat{A}_t = R_t - b(s_t)$.

 Re-fit the baseline, by minimizing $\|b(s_t) - R_t\|^2$,
 summed over all trajectories and timesteps.


 Update the policy, using a policy gradient estimate \hat{g} ,
 which is a sum of terms $\nabla_{\theta} \log \pi(a_t | s_t, \theta) \hat{A}_t$

end for


- ▶ A \hat{V}_{ϕ} or \hat{Q}_{ϕ} baseline provides a value even in unseen states
- ▶ Recompute the baseline from all trajectories
- ▶ Or update the baseline from one trajectory
- ▶ If the critic is estimated based on the previous critic, it becomes bootstrap

https://www.youtube.com/watch?v=S_gwYj1Q-44 (36')

Monte Carlo versus Bootstrap approaches



Monte Carlo approach



Bootstrap approach

- ▶ Trajectory-based approach: Monte Carlo methods
- ▶ Does not record a critic from an iteration to the next
- ▶ Gets an unbiased estimate for all visited state-action pairs using the current batch
- ▶ Bootstrap approaches: record a parametrized critic
- ▶ Bootstrap is sample efficient but suffers from bias and is unstable
- ▶ Monte Carlo is stable, but suffers from variance and is slower

Estimating $V^\pi(s)$ or $Q^\pi(s, a)$

- ▶ Let us define \hat{V}_ϕ or \hat{Q}_ϕ as estimators of $V^\pi(s)$ or $Q^\pi(s, a)$
- ▶ Two approaches to estimate them:
 - ▶ **Monte Carlo estimate:** Regression against empirical return

$$\phi_{j+1} \rightarrow \arg \min_{\phi_j} \frac{1}{m} \sum_{i=1}^m \sum_{t=1}^H (\hat{V}_{\phi_j}^\pi(s_t^{(i)}) - (\sum_{k=t}^H r(s_t^{(i)}, a_t^{(i)})))^2$$

- ▶ **Temporal difference estimate:** init $\hat{V}_{\phi_0}^\pi$ and fit using (s, a, r, s') data

$$\phi_{j+1} \rightarrow \min_{\phi_j} \sum_{(s, a, r, s')} \|r + \gamma \hat{V}_{\phi_j}^\pi(s') - \hat{V}_{\phi_j}^\pi(s)\|^2$$

May need some regularization to prevent large steps in ϕ

- ▶ Similar equations for \hat{Q}_ϕ



Martin Riedmiller. Neural fitted q iteration—first experiences with a data efficient neural reinforcement learning method. In *European Conference on Machine Learning*, pp. 317–328. Springer, 2005



András Antos, Csaba Szepesvári, and Rémi Munos. Fitted Q-iteration in continuous action-space MDPs. In *Advances in neural information processing systems*, pp.9–16, 2008.


Being truly actor-critic

- ▶ “Although the REINFORCE-with-baseline method learns both a policy and a state-value function, we do not consider it to be an actor–critic method because its state-value function is used only as a baseline, not as a critic.”
- ▶ “That is, it is not used for bootstrapping (updating the value estimate for a state from the estimated values of subsequent states), but only as a baseline for the state whose estimate is being updated.”
- ▶ “This is a useful distinction, for only through bootstrapping do we introduce bias and an asymptotic dependence on the quality of the function approximation.”



Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction (Second edition)*. MIT Press, 2018, p. 331

Practical implementation of critics



- \hat{V}_ϕ is smaller, but not necessarily easier to estimate

Synthesis

$$\nabla_{\theta} J(\theta) = \mathbb{E}[\psi_t \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)})]$$

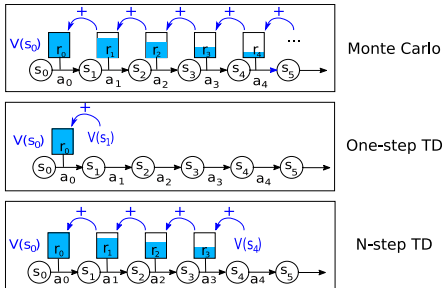
where ψ_t can be:

1. $\psi_t = \sum_{t'=0}^H \gamma^{t'} r_{t'}$: total reward of trajectory
2. $\psi_t = \sum_{t'=t}^H \gamma^{t'-t} r_{t'}$: sum of rewards after a_t
3. $\psi_t = \sum_{t'=t}^H \gamma^{t'-t} r_{t'} - b(s_t)$: sum of rewards after a_t with baseline
4. $\psi_t = r_t + \gamma V^{\pi}(s_{t+1}) - V^{\pi}(s_t) = \delta_t$: TD error, with
 $V^{\pi}(s_t) = \mathbb{E}_{a_t} [\sum_{l=0}^H \gamma^l r_{t+l}]$
5. $\psi_t = Q^{\pi}(s_t, a_t)$: state-action value function, with
 $Q^{\pi}(s_t, a_t) = \mathbb{E}_{a_{t+1}} [\sum_{l=0}^H \gamma^l r_{t+l}]$
6. $\psi_t = A^{\pi}(s_t, a_t)$: **advantage function**, with
 $A^{\pi}(s_t, a_t) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t) = \mathbb{E}[\delta_t]$




John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. *arXiv preprint arXiv:1506.02438*, 2015

Monte Carlo, One-step TD and N-step return



- ▶ One-step TD suffers from bias
- ▶ MC suffers from variance due to exploration (+ stochastic trajectories)
- ▶ MC is on-policy → less sample efficient
- ▶ N-step TD: tuning N to control the bias-variance compromise

Bias-variance compromise



- ▶ Total error = $\text{bias}^2 + \text{variance} + \text{irreducible error}$
- ▶ A more complex model (e.g. bigger network) generally has more variance, but less bias
- ▶ Tuning N in the N -step return helps finding the right compromise.

Generalized Advantage Estimation: λ return

- ▶ The N-step return can be reformulated using a continuous parameter λ
- ▶ $\hat{A}_{\phi}^{(\gamma, \lambda)} = \sum_{l=0}^H (\gamma \lambda)^l \delta_{t+l}$
- ▶ $\hat{A}_{\phi}^{(\gamma, 0)} = \delta_t = \text{one-step return}$
- ▶ $\hat{A}_{\phi}^{(\gamma, 1)} = \sum_{l=0}^H (\gamma)^l \delta_{t+l} = \text{MC estimate}$
- ▶ The λ return comes from eligibility trace methods
- ▶ Provides a continuous grip on the bias-variance trade-off



John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. *arXiv preprint arXiv:1506.02438*, 2015

Computing \hat{V}_ϕ or \hat{Q}_ϕ One-step, N-step, or λ return

Monte Carlo:
no bias, higher
variance, lower
sample efficiency

bootstrap:
bias, lower
variance, higher
sample efficiency

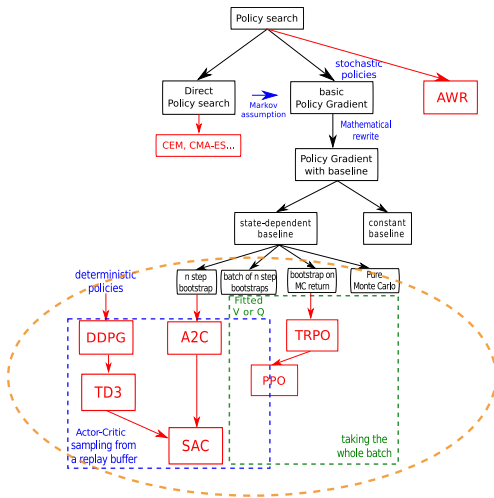
Batch Monte Carlo estimate TRPO	Incremental Monte Carlo estimate PPO
Batch Temporal Difference estimate	Incremental Temporal Difference estimate (actor-critic) A2C, SAC

- ▶ There are many possibilities to approximate the policy gradient
 - ▶ Six proxies to advantage estimators
 - ▶ Batch, N-step return, λ return, one-step TD update...
- ▶ Results in a large variety of algorithms

Take home messages: DPS vs PG

- ▶ Direct policy search:
 - ▶ optimization without a utility model
 - ▶ derivative-free
 - ▶ poor sample reuse, low sample efficiency
- ▶ Policy Gradient:
 - ▶ Uses analytical derivative of the policy function
 - ▶ Uses information from each step, not just trajectories
 - ▶ A baseline is used to reduce variance
 - ▶ When bootstrap comes into play, becomes actor-critic

Global view



- A second video will focus on state-of-the-art algorithms

Any question?



Send mail to: Olivier.Sigaud@upmc.fr



András Antos, Csaba Szepesvári, and Rémi Munos.

Fitted q-iteration in continuous action-space mdps.

In *Advances in neural information processing systems*, pp. 9–16, 2008.



Marc Peter Deisenroth, Gerhard Neumann, Jan Peters, et al.

A survey on policy search for robotics.

Foundations and Trends® in Robotics, 2(1–2):1–142, 2013.



Martin Riedmiller.

Neural fitted q iteration—first experiences with a data efficient neural reinforcement learning method.

In *European Conference on Machine Learning*, pp. 317–328. Springer, 2005.



John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel.

High-dimensional continuous control using generalized advantage estimation.

arXiv preprint arXiv:1506.02438, 2015.



Olivier Sigaud and Freek Stulp.

Policy search in continuous action domains: an overview.

Neural Networks, 113:28–40, 2019.



Richard S. Sutton and Andrew G. Barto.

Reinforcement Learning: An Introduction.

MIT Press, 1998.



Richard S. Sutton and Andrew G. Barto.

Reinforcement Learning: An Introduction (Second edition).

MIT Press, 2018.



Ronald J. Williams.

Simple statistical gradient-following algorithms for connectionist reinforcement learning.

Machine Learning, 8(3-4):229–256, May 1992.

ISSN 0885-6125.

doi: 10.1007/BF00992696.